



# 0. Course Details & Introduction

COMP7404

Computational Intelligence and Machine Learning

Dirk Schnieders

# Course Staff

- Instructors
  - Dirk Schnieders
    - Email: [sdirk@cs.hku.hk](mailto:sdirk@cs.hku.hk)
    - Consultation hours: Tue 2:30pm - 3:30pm & by appointment
- TA
  - Zhao Shihao
    - Email: [shzhao@connect.hku.hk](mailto:shzhao@connect.hku.hk)
    - Consultation hours: Tue 5:00pm - 6:00pm & by appointment

# Schedule

#	Date	Venue	Time	Topic*
1	2 Sep	CYP-P4	7pm - 10pm	Course Details / Uninformed Search
2	9 Sep	CYP-P4	7pm - 10pm	Informed Search
3	16 Sep	CYP-P4	7pm - 10pm	Local Search CSP
4	23 Sep	CYP-P4	7pm - 10pm	Markov Decision Process
5	30 Sep	CYP-P4	7pm - 10pm	Reinforcement Learning
6	7 Oct	CYP-P4	7pm - 10pm	Perceptron & Adaline
7	21 Oct	CYP-P4	7pm - 10pm	Logistic Regression, Support Vector Machines, Decision Tree Learning, KNN
8	28 Oct	CYP-P4	7pm - 10pm	Evaluation / Tuning & Ensemble Learning
9	4 Nov	CYP-P4	7pm - 10pm	Presentations
10	11 Nov	CYP-P4	7pm - 10pm	Presentations

# Assessment

- Assignments (30 %)
  - A1 (8%): Search (13 Sep - 27 Sep)
  - A2 (8%): Adversarial Search (27 Sep - 11 Oct)
  - A3 (8%): Reinforcement Learning (25 Oct - 8 Nov)
  - A4 (6%): Machine Learning (11 Nov - 25 Nov)
- Group project (20 %)
  - Read and implement machine learning research paper
  - Presentation to class (8 minutes presentation)
  - Max 36 Groups
  - Max 4 students per group
- Final examination (50 %): written, closed book, 2 hours

# Prerequisite

- Python programming
  - Students without strong experience may spend a lot of time on assignments
- Basic linear algebra, calculus and statistics

# Assessment - Project - Group Selection

- Select your group on Moodle
  - Start: 16 Sep
  - Deadline: 23 Sep 23:59
  - We will randomly assign a group for you if you don't make a selection by the deadline
- Member swapping
  - Allowed if approved before start of project
    - All group members of both groups must agree

# Assessment - Project - Paper Selection

- Read a machine learning paper
- Implement / test the paper
  - TensorFlow 2 / PyTorch on CS [GPU Farm](#)
  - Employ own test/training data
  - Make your code available to others
- List of project topics will be made available on 28 Sep
- Topic selection
  - On 7 Oct in-class
    - First come first serve
    - Done by group leader
  - Deadline: 8 Oct 23:59

# Plagiarism

- What is Plagiarism ?
  - <https://tl.hku.hk/plagiarism/>
- If a student commits plagiarism, with evidence after investigation, no matter whether the student concerned admits or not, a penalty will be imposed
- First Attempt: the student shall be warned in writing and receive zero mark for the whole assignment or the whole test; if the student does not agree, s/he can appeal to the Programme Director within a week
- Subsequent Attempt: May impose any of the following penalties: a published reprimand, suspension of study for a period of time, fine, or expulsion from the University
- Both the student who copies and the student who offers his/her work for copying shall be penalized



# Course Materials

- Available on Moodle
  - Use the provided materials responsibly
  - For your own research and private study only
  - Don't distributed to others without the appropriate authorization
    - We do not hold copyright for most materials

# References and Acknowledgements

- Textbooks
  - Artificial Intelligence: A Modern Approach, Third Edition, Stuart Russell & Peter Norvig
  - Python Machine Learning (PML), Third Edition Sebastian Raschka & Vahid Mirjalili
  - Introduction to Machine Learning (IML), Third Edition Ethem Alpaydin
- Other Courses
  - [CS188](#), Berkeley
  - [Machine Learning Crash Course](#), Google
  - [Machine Learning](#), Coursera
- Additional references (if any) will be listed on the lecture slides

# Tools

- Provided examples and exercises will use the following tools
  - [Python 3](#)
  - [NumPy](#)
  - [SciPy library](#)
  - [scikit-learn](#)
  - [Matplotlib](#)
  - [Pandas](#)
  - [Tensorflow](#)
  - [Jupyter Notebook](#)
- All of the above can be installed conveniently with [Anaconda](#)
  - We recommend to use [Google CoLab](#)

# Discussion Forum

- If you have questions about course materials, please **use the forum**
  - Answer questions from other students!

Q & A

# Introduction

# Intelligence

- We call ourselves Homo sapiens
  - Latin: Wise Man
  - Intelligence is important to us
- For thousands of years, we have tried to understand how we humans think
- Intelligence is most widely studied in humans, but has also been observed in animals and in plants

# Artificial Intelligence

- AI goes further than just understanding intelligence
  - It attempts to build intelligent entities
- Is AI science, or is it engineering?
  - AI's science goal
    - To understand the principles and mechanism that account for intelligent action
  - AI's engineering goal
    - To design intelligent artifacts that can survive and operate in the physical world and solve problems of considerable scientific difficulty at high levels of competence



# Goal of AI

- Should we create machines that ...
  - A. Think Humanly
  - B. Act Humanly
  - C. Think Rationally
  - D. Act Rationally





# A. Think Humanly

- Researchers find out how we think by
  - Introspection
  - Psychological experiments
  - Brain imaging
- Cognitive science constructs precise and testable theories of the human mind
  - E.g., express a theory as a computer program and compare input-output behaviors to a human
  - If there is a match, some of the programs mechanism could also be operating in humans



# A. Think Humanly

- The human brain is one of the great mysteries of science
  - How does our brain process information?
- The brain consists of nerve cells (aka neurons) and the collection of these simple cells leads to thought, action and consciousness
- The recent development of functional magnetic resonance imaging (fMRI) provides neuroscientists with details of brain activities



# A. Think Humanly

- Brains and digital computers have somewhat different properties
- A crude comparison of the raw computational resources

	Supercomputer	Personal Computer	Human Brain
Computational units	$10^6$ GPUs + CPUs	8 CPU cores	$10^6$ columns
	$10^{15}$ transistors	$10^{10}$ transistors	$10^{11}$ neurons
Storage units	$10^{16}$ bytes RAM	$10^{10}$ bytes RAM	$10^{11}$ neurons
	$10^{17}$ bytes disk	$10^{12}$ bytes disk	$10^{14}$ synapses
Cycle time	$10^{-9}$ sec	$10^{-9}$ sec	$10^{-3}$ sec
Operations/sec	$10^{18}$	$10^{10}$	$10^{17}$

- Would we be able to achieve the brain's level of intelligence with a computer of unlimited capacity?

# Should we create machines that ...

- A. Think Humanly
- B. Act Humanly
- C. Think Rationally
- D. Act Rationally


## B. Act Humanly



- The Turing Test (aka Imitation Game) was designed to provide a definition of intelligence
- A computer passes the test if a human interrogator, after posing some questions, cannot tell whether the response come from a human or a computer



Human  
interrogator




## B. Act Humanly

- The underlying principles of intelligence are more important than to duplicate an exemplar
- Consider another field: Artificial Flight
  - The Wright brothers succeeded because they stopped imitating birds and started using wind tunnels and learn about aerodynamics
  - It was not their goal to make “machines that fly so exactly like pigeons that they can fool even other pigeons”



## B. Act Humanly

- A reverse Turing test is a Turing test in which the objective or roles between computers and humans have been reversed. Interrogator is a computer. Interrogatee is a human.
  - Example: CAPTCHA
    - It is program that can generate and grade tests that
      - most humans can pass, but
      - current computer programs cannot pass





# C. Think Rationally

- What are the laws that guide and underlie our thinking?
- Greek schools developed various forms of logic
  - Notation and rules of derivation for thoughts
  - Example: Socrates is a man; all men are mortal; therefore, Socrates is mortal
- By 1965, programs existed that could (in principle) solve any solvable problem described in logic notation
- Problems with this approach
  - How to take informal knowledge and state it in formal terms?  
How about uncertainty?

# D. Act Rationally

- Act so as to achieve the best outcome or, when there is uncertainty, the best expected outcome
- Advantages over the other approaches
  - More general than the laws of thought approach because correct inference is just one of several possible mechanisms for achieving rationality
  - The standard of rationality is mathematically well defined. Human behavior, on the other hand, is well adapted only for one specific environment
- In this course we will focus on the general principles of rational agents and how to build them
  - An agent is something that perceives and acts

# Rational Agent

- This course is about designing rational agents
- For any given class of environments and tasks, we seek the agent (or class of agents) with the best performance
  - Computational limitations make perfect rationality unachievable
  - Design best program for given machine resources

# Machine Learning

- An agent is learning if it improves its performance on future tasks
- Why would we want an agent to learn? If the design of an agent can be improved, why not design the agent with that improvement to begin with?
  - Cannot anticipate all possible situations
  - Cannot anticipate all changes over time
  - Don't know how to program some solutions

# History of AI



# History of AI - Turing Award Winners

- Marvin Minsky (1969)
- John McCarthy (1971)
- Edward Feigenbaum and Raj Reddy (1994)
- Judea Pearl (2011)
- Yoshua Bengio, Geoffrey Hinton, and Yann LeCun (2019)

# History of AI - Milestones

- Inception (1943 - 1956)
- Early Enthusiasm (1952 - 1969)
- A dose of reality (1966 - 1973)
- Expert systems (1969 - 1986)
- Return of NN (1986 - present)
- Probabilistic reasoning (1987 - present)
- Big data (2001 - present)
- Deep Learning (2011 - present)

# The State of the Art

- Publications
  - AI papers increased 20 fold between 2010 to 2019 to 20,000 a year
- Conferences
  - Attendance of [NeurIPS](#) increased 800% since 2012 to 13,500
- Industry
  - AI startups in the US increased 20 fold to over 800 from 2010 to 2019
- Internationalization
  - China publishes more AI papers per year then US and about as many as Europe
  - In citation weighted impact, US is ahead by 50% vs. China



# The State of the Art

- Vision
  - Error rates for object detection improved from 28% to less than 2%
- Speed
  - Training time for image recognition dropped by a factor of 100 in last 2 years
  - Amount of computing power used in top AI applications is doubling every few month
- Humans vs. AI
  - AI is better in chess, go, poker, pac-man, jeopardy!, object detection, speech recognition in limited domain, chinese-to-english in restricted domain, Quake III, Dota 2, StarCraft II, many Atari games, Skin cancer detection, prostate cancer detection, protein folding, ...

# Benefits of AI

First solve AI, then use AI to solve everything else.

Demis Hassabis, Google DeepMind



# Risks of AI

- Lethal autonomous weapons
- Surveillance
- Biased decision making
- Impact on employment
- Safety-critical applications
- Cybersecurity

# Risks of AI - Superhuman AI

- Most experts agree that we will eventually be able to create a superhuman AI
  - An intelligence that far surpasses human ability

# Risks of AI - The Gorilla Problem

- About seven million years ago, a now-extinct primate evolved
  - one branch led to gorillas
  - another to humans
- Today the gorillas are probably not too happy about the human branch
  - They have no control over their future



# Risks of AI - The Gorilla Problem

- If the gorilla problem is the result of developing AI then we should stop working on it
- If superhuman AI were a black box from outer space, we should be careful in opening the box
  - But it is not, **we** design the AI systems
  - If AI does end up taking control, it would be a design failure
- We need to understand the source of potential failure
  - Philosophical foundations of AI
  - Maybe the most important area of AI research